

Artificial Intelligence

Neural Networks

Lesson 12: Learning Vector Quantization

Vincenzo Piuri

Università degli Studi di Milano

Contents



- Learning vector quantization
- Learning vector quantization networks
- Learning rules
- Soft learning vector quantization
- Expectation maximization

Learning Vector Quantization (1)

- So far: fixed learning tasks
 - The data consists of input/output pairs
 - The objective is to produce desired output for given input
 - This allows to describe training as error minimization
- Now: **free learning tasks**
 - The data consists only of input values/vectors
 - The objective is to produce similar output for similar input (clustering)

Learning Vector Quantization (2)

- **Delaunay Triangulation:** simple triangle (shown in gray on the left)
- **Voronoi Diagram:** mid-perpendiculars of the triangle's edges (shown in blue on the left, in gray on the right)




Learning Vector Quantization (3)

- **Learning Vector Quantization**

- Find a suitable quantization (many-to-few mapping, often to a finite set) of the input space, e.g. a tessellation of a Euclidean space
- Training adapts the coordinates of so-called reference or codebook vectors, each of which defines a region in the input space

Learning Vector Quantization (4)

- Finding clusters in a given set of data points
 - Data points are represented by empty circles (\circ)
 - Cluster centers are represented by full circles (\bullet)



Learning Vector Quantization Nets (1)

- A **Learning Vector Quantization Network (LVQ)** is a feed-forward 2-layered neural network
- It can be viewed as a RBF network with hidden layer used as output layer
- The network input function of each output neuron is a distance function of the input vector and the weight vector

- $\forall u \in U_{out}: f_{net}^{(u)}(\vec{w}_u, \vec{in}_u) = d(\vec{w}_u, \vec{in}_u)$

- $d(\vec{x}, \vec{y}) = 0 \iff \vec{x} = \vec{y}$


- $d(\vec{x}, \vec{y}) = d(\vec{y}, \vec{x})$ (Symmetry)

- $d(\vec{x}, \vec{z}) \leq d(\vec{x}, \vec{y}) + d(\vec{y}, \vec{z})$ (Triangle inequality)

Learning Vector Quantization Nets (2)

- Distance functions: **Minkowski Family**

- $d_k(\vec{x}, \vec{y}) = (\sum_{i=1}^n |x_i - y_i|^k)^{\frac{1}{k}}$
- $k = 1$: Manhattan or city block distance
- $k = 2$: Euclidean distance
- ...
- $k \rightarrow \infty$: Maximum distance



Learning Vector Quantization Nets (3)

- The activation function of each output neuron is a **radial function**
 - Monotonically decreasing function
 - $f: \mathbb{R}_0^+ \rightarrow [0,1]$ with $f(0) = 1$ and $\lim_{x \rightarrow \infty} f(x) = 0$

Learning Vector Quantization Nets (4)


- The output function of each output neuron is not a simple function of the activation of the neuron
 - It considers the activations of all output neurons
 - $f_{out}^{(u)}(act_u) = \begin{cases} 1 & \text{if } act_u = \max_{v \in U_{out}} act_v \\ 0 & \text{otherwise} \end{cases}$
 - If more than one unit has the maximal activation, one is selected at random to have an output of 1, all others are set to output 0: **winner-takes-all principle**

Learning Vector Quantization Nets (5)

- Radial activation functions


rectangle function:

$$f_{\text{act}}(\text{net}, \sigma) = \begin{cases} 0, & \text{if } \text{net} > \sigma, \\ 1, & \text{otherwise.} \end{cases}$$




triangle function:

$$f_{\text{act}}(\text{net}, \sigma) = \begin{cases} 0, & \text{if } \text{net} > \sigma, \\ 1 - \frac{\text{net}}{\sigma}, & \text{otherwise.} \end{cases}$$




cosine until zero:

$$f_{\text{act}}(\text{net}, \sigma) = \begin{cases} 0, & \text{if } \text{net} > 2\sigma, \\ \frac{\cos(\frac{\pi}{2\sigma} \text{net}) + 1}{2}, & \text{otherwise.} \end{cases}$$



Gaussian function:

$$f_{\text{act}}(\text{net}, \sigma) = e^{-\frac{\text{net}^2}{2\sigma^2}}$$





Learning rules (1)

- Adaptation of reference vectors
 - For each training pattern find the closest reference vector
 - Adapt only this reference vector (winner neuron)
 - For classified data the class may be considered
 - Each reference vector is assigned to a class

Learning rules (2)


- **Attraction rule** (data point and reference vector have same class)
 - $\vec{r}^{(new)} = \vec{r}^{(old)} + \eta(\vec{x} - \vec{r}^{(old)})$
- **Repulsion rule** (data point and reference vector have different class)
 - $\vec{r}^{(new)} = \vec{r}^{(old)} - \eta(\vec{x} - \vec{r}^{(old)})$

Learning rules (3)




- Adaptation of reference vectors

Learning rules (4)





Online training



Batch training

- Adaptation of reference vectors

Learning rules (5)



- Fixed learning rate can lead to oscillations
 - Solution: **time dependent learning rate**

Learning rules (6)

- Update rule for **classified data**
 - Update not only the one reference vector that is closest to the data point (the winner neuron), but update the two closest reference vectors
 - All other reference vectors remain unchanged

Learning rules (7)

- Standard learning vector quantization may drive the reference vectors further and further apart
- **Window rule**
 - Update only if the data point \vec{x} is close to the classification boundary

$$\min \left(\frac{d(\vec{x}, \vec{r}_j)}{d(\vec{x}, \vec{r}_k)}, \frac{d(\vec{x}, \vec{r}_k)}{d(\vec{x}, \vec{r}_j)} \right) > \theta, \quad \text{where} \quad \theta = \frac{1 - \xi}{1 + \xi}.$$

- ξ is a parameter that has to be specified by a user
- ξ describes the “width” of the window around the classification boundary
- The update ceases once the classification boundary has been moved far enough away

Soft Learning Vector Quantization

- Use soft assignments instead of winner-takes-all
 - Assumption: given data was sampled from a mixture of normal distributions
- Closely related to clustering by estimating a **mixture of Gaussians**
 - (Crisp or hard) learning vector quantization is an “online version” of C-means clustering
 - Soft learning vector quantization is an “online version” of estimating a mixture of Gaussians

Expectation Maximization (1)

- Mixture of Gaussians

- Assumption: Data was generated by sampling a set of normal distributions
- We assume that the probability density can be described as

$$f_{\vec{X}}(\vec{x}; \mathbf{C}) = \sum_{y=1}^c f_{\vec{X}, Y}(\vec{x}, y; \mathbf{C}) = \sum_{y=1}^c p_Y(y; \mathbf{C}) \cdot f_{\vec{X}|Y}(\vec{x}|y; \mathbf{C}).$$

\mathbf{C}	is the set of cluster parameters
\vec{X}	is a random vector that has the data space as its domain
Y	is a random variable that has the cluster indices as possible values (i.e., $\text{dom}(\vec{X}) = \mathbb{R}^m$ and $\text{dom}(Y) = \{1, \dots, c\}$)
$p_Y(y; \mathbf{C})$	is the probability that a data point belongs to (is generated by) the y -th component of the mixture
$f_{\vec{X} Y}(\vec{x} y; \mathbf{C})$	is the conditional probability density function of a data point given the cluster (specified by the cluster index y)

Expectation Maximization (2)

- Maximum likelihood estimation of the cluster parameters

- The likelihood function is difficult to optimize

$$L(\mathbf{X}; \mathbf{C}) = \prod_{j=1}^n f_{\vec{X}_j}(\vec{x}_j; \mathbf{C}) = \prod_{j=1}^n \sum_{y=1}^c p_Y(y; \mathbf{C}) \cdot f_{\vec{X}|Y}(\vec{x}_j|y; \mathbf{C}),$$

- Approach: Assume “hidden” variables Y_j stating the clusters that generated the data points \vec{x}_j

$$L(\mathbf{X}, \vec{y}; \mathbf{C}) = \prod_{j=1}^n f_{\vec{X}_j, Y_j}(\vec{x}_j, y_j; \mathbf{C}) = \prod_{j=1}^n p_{Y_j}(y_j; \mathbf{C}) \cdot f_{\vec{X}_j|Y_j}(\vec{x}_j|y_j; \mathbf{C}).$$

- Problem: We do not know the values of Y_j

Expectation Maximization (3)

- Approach

- See the Y_j as random variables (the values y_j are not fixed) and consider a probability distribution over the possible values
- $L(\mathbf{X}, \vec{y}; \mathbf{C})$ becomes a random variable, even for a fixed data set \mathbf{X} and fixed cluster parameters \mathbf{C}
- Try to maximize the expected value of $L(\mathbf{X}, \vec{y}; \mathbf{C})$ or $\ln L(\mathbf{X}, \vec{y}; \mathbf{C})$ (expectation maximization)

Expectation Maximization (4)

- Find the cluster parameters

$$\hat{\mathbf{C}} = \underset{\mathbf{C}}{\operatorname{argmax}} E([\ln]L(\mathbf{X}, \vec{y}; \mathbf{C}) \mid \mathbf{X}; \mathbf{C}),$$

- Maximize the expected likelihood

$$E(L(\mathbf{X}, \vec{y}; \mathbf{C}) \mid \mathbf{X}; \mathbf{C}) = \sum_{\vec{y} \in \{1, \dots, c\}^n} p_{\vec{Y}|\mathcal{X}}(\vec{y}|\mathbf{X}; \mathbf{C}) \cdot \prod_{j=1}^n f_{\vec{X}_j, Y_j}(\vec{x}_j, y_j; \mathbf{C})$$

- Or maximize the expected log-likelihood

$$E(\ln L(\mathbf{X}, \vec{y}; \mathbf{C}) \mid \mathbf{X}; \mathbf{C}) = \sum_{\vec{y} \in \{1, \dots, c\}^n} p_{\vec{Y}|\mathcal{X}}(\vec{y}|\mathbf{X}; \mathbf{C}) \cdot \sum_{j=1}^n \ln f_{\vec{X}_j, Y_j}(\vec{x}_j, y_j; \mathbf{C}).$$

Expectation Maximization (5)

- Still difficult to optimize directly
 - Solution: use the equation as an iterative scheme, fixing C in some terms
 - Iteratively compute better approximations

Expectation Maximization (6)

- Iterative scheme for expectation maximization
 - Choose some initial set C_0 of cluster parameters and then compute

$$\begin{aligned} C_{k+1} &= \operatorname{argmax}_{\mathbf{C}} E(\ln L(\mathbf{X}, \vec{y}; \mathbf{C}) \mid \mathbf{X}; \mathbf{C}_k) \\ &= \operatorname{argmax}_{\mathbf{C}} \sum_{\vec{y} \in \{1, \dots, c\}^n} p_{\vec{Y}|\mathcal{X}}(\vec{y}|\mathbf{X}; \mathbf{C}_k) \sum_{j=1}^n \ln f_{\vec{X}_j, Y_j}(\vec{x}_j, y_j; \mathbf{C}) \\ &= \operatorname{argmax}_{\mathbf{C}} \sum_{\vec{y} \in \{1, \dots, c\}^n} \left(\prod_{l=1}^n p_{Y_l|\vec{X}_l}(y_l|\vec{x}_l; \mathbf{C}_k) \right) \sum_{j=1}^n \ln f_{\vec{X}_j, Y_j}(\vec{x}_j, y_j; \mathbf{C}) \\ &= \operatorname{argmax}_{\mathbf{C}} \sum_{i=1}^c \sum_{j=1}^n p_{Y_j|\vec{X}_j}(i|\vec{x}_j; \mathbf{C}_k) \cdot \ln f_{\vec{X}_j, Y_j}(\vec{x}_j, i; \mathbf{C}). \end{aligned}$$

- Each EM iteration increases the likelihood of the data and the algorithm converges to a local maximum of the likelihood function

Expectation Maximization (7)

- Core Iteration Formula

$$\mathbf{C}_{k+1} = \underset{\mathbf{C}}{\operatorname{argmax}} \sum_{i=1}^c \sum_{j=1}^n p_{Y_j|\vec{X}_j}(i|\vec{x}_j; \mathbf{C}_k) \cdot \ln f_{\vec{X}_j, Y_j}(\vec{x}_j, i; \mathbf{C})$$

1. Expectation step

- For all data points \vec{x}_j : Compute for each normal distribution the probability $p_{Y_j|\vec{X}_j}(i|\vec{x}_j; \mathbf{C}_k)$ that the data point was generated from it

2. Maximization step

- For all normal distributions: estimate the parameters by standard maximum likelihood estimation using the probabilities (“weights”) assigned to the data points w.r.t. the distribution in the expectation step